

1

2

3

4

5

6

7

8

9

10

<

1 of 40

>

⌚ All

Question # 1

✎ Revisit

Why are big data applications liable to latency?

Choose the best option

- ☐ Big data cannot use in-memory computing
- ☐ Big data applications are still in the early stages of development.
- ☐ The volume of big data is too large to be analyzed rapidly
- ☐ Big data may reside in a different location from the application

Question # 5

 Revisit

Choose the best option

- ☐ Moving data to computation
- ☐ Moving computation to data
- ☐ Moving data to computation and Moving computation to data
- ☐ None of the above

Data Locality means:

Question # 4

 Revisit

Which of the following is fundamental datastructure of Spark?

Choose the best option

- ☐ RDD
- ☐ Dataframe
- ☐ Dataset
- ☐ None of the above

Question # 3

Typically a Hbase Region server is collocated with

 Revisit

Choose the best option

- ☐ HDFS namenode
- ☐ HDFS datanode
- ☐ As a client to HDFS server
- ☐ Resource manager

1

2

3

4

5

6

7

8

9

10

<

2 of 40

>

All



Question # 2

Command is used to know the status of all the daemons of hadoop.

 Revisit

Choose the best option

- ☐ fsck
- ☐ distcp
- ☐ jps
- ☐ None of the above

Question # 8

 Revisit

Choose the best option

- ☐ GraphX
- ☐ MLlib
- ☐ Spark SQL
- ☐ Spark R

Which of the following is module for Structured data processing?

Question # 7

 Revisit

Choose the best option

What message is generated by a datanode to indicate its connectivity with namenode?

- ☐ Beap
- ☐ Heartbeat
- ☐ Analog Pulse
- ☐ Map

Question # 6

 Revisit

Choose the best option

_____ is not a component of spark.

- ☐ SparkR
- ☐ Mlib
- ☐ GraphX
- ☐ Sqoop

11

12

13

14

15

16

17

18

19

20

<

17 of 40

>

All

Question # 17

 Revisit

Pick the correct statement about hadoop when compared to RDBMS?

Choose the best option

- ☐ Hadoop is suitable for read and write many times
- ☐ Hadoop does ACID transactions
- ☐ Hadoop works better on unstructured and semi-structured
- ☐ Hadoop has higher data Integrity



Question # 15

 Revisit

Why do we use the SSH in Hadoop Cluster?

Choose the best option

- ☐ To perform the Passwordless authentication
- ☐ To establish the communication between Master
- ☐ Both of the above
- ☐ None of The above

16

17

18

19

20

21

22

23

24

25

<

21 of 40

>

All

Question # 21

 Revisit

Which of the following services is provided by YARN?

Choose the best option

- ☐ Global resource management
- ☐ Record reader
- ☐ MapReduce engine
- ☒ Data Mining

11

12

13

14

15

16

17

18

19

20

<

19 of 40

>

All

Question # 19

 Revisit

Hive provides a SQL like language called

Choose the best option

☐ SQL Hive

☐ Hive QL

☐ DB QL

☐ Hive Data

Question # 23

 Revisit

Pick the false statement about hadoop:

Choose the best option

- ☐ The main algorithm used in it is Map Reduce
- ☐ It is a distributed framework
- ☐ It runs with commodity hard ware
- ☐ All are true



21

22

23

24

25

26

27

28

29

30

<

29 of 40

>

All

Question # 29

 Revisit

Which of the following is not true?

Choose the best option

- ☐ In Pseudo distributed mode no daemons will be running
- ☐ In fully distributed mode all the daemons will be running
- ☐ In standalone no daemons will be running
- ☐ In Pseudo distributed mode no daemons will be running
daemons will be running on same machine

21

22

23

24

25

26

27

28

29

30

<

27 of 40

>

All

Question # 27

 Revisit

The main advantage of creating table partition is

Choose the best option

- ☐ Effective storage memory utilization
- ☐ faster query performance
- ☐ Less RAM required by namenode
- ☐ simpler query syntax



21

22

23

24

25

26

27

28

29

30

<

28 of 40

>

All

Question # 28

 Revisit

What can best be described as a programming model used to develop Hadoop-based applications that can process massive amounts of data?

Choose the best option

- ☐ Oozie
- ☐ Mahout
- ☐ MapReduce engine
- ☐ All of the mentioned

Question # 29

🔄 Revisit

Which of the following is not true?

Choose the best option

- ☐ In Pseudo distributed mode no daemons will be running
- ☐ In fully distributed mode all the daemons will be running on same machine
- ☐ In standalone no daemons will be running
- ☐ In Pseudo distributed mode no daemons will be running and In fully distributed daemons will be running on same machine

Question # 33

 Revisit

Which property gets configured in mapred-site.xml?

Choose the best option

- ☐ Host and port where MapReduce job runs.
- ☐ Java Environment variables.
- ☐ Directory names to store hdfs files
- ☐ Replication Factor

Question # 35

 Revisit

On dropping a managed/internal table

Choose the best option

- ☐ The schema gets dropped without dropping the data
- ☐ The data gets dropped without dropping the schema
- ☐ An error is thrown
- ☐ Both the schema and the data is dropped

Question # 38

 Revisit

To retrieve all rows of a table in HBase, we use:

Choose the best option

- ☐ get
- ☐ scan
- ☐ put
- ☐ select

31

32

33

34

35

36

37

38

39

40

<

37 of 40

>

All

Question # 37


 Revisit

What role does the map function play in a word count query?

Choose the best option

- ☐ It sorts the words alphabetically and returns a list of the most frequent words.
- ☐ It creates a list with each word as a key and the number of occurrences as a value.
- ☐ It creates a list with each word as a key and every occurrence as a value.
- ☐ It returns a list with each document as a key and the number of words as a value.

Question # 40



We have a 500 MB file and HDFS block size of 64 MB.
We are using `hadoop fs -put` command to write this file. Just after this command has finished writing 400 MB of this file, what would another user see when trying to access this file?

Choose the best option

- ☐ They would see no content until the whole file written and closed.
- ☐ They would see the current state of the file through the last completed block
- ☐ They would see the current state of the file, up to the last bit written by the command
- ☐ They would see Hadoop throw a `ConcurrentFileAccessException` when they try to access this file.

Question # 39

 Revisit

Which of the following deals with small files issue?

Choose the best option

- ☐ Sequence files
- ☐ Hadoop Archives
- ☐ HBase
- ☐ All of the above

Question # 34

 Revisit

Apache HBase was modeled after Google's _____

Choose the best option

- ☐ Foundation DB
- ☐ Big top
- ☐ Big Tables
- ☐ None of the above

26

27

28

29

30

31

32

33

34

35

<

35 of 40

>

All

Question # 35

 Revisit

On dropping a managed/internal table

Choose the best option

- ☐ The schema gets dropped without dropping the data
- ☐ The data gets dropped without dropping the schema
- ☐ An error is thrown
- ☒ Both the schema and the data is dropped

Clear Response

Question # 36

 Revisit

If a Big data analyst were to analyze data from a database of call logs provided by a telecom service provider, which element of big data would be dealing with?

Choose the best option

☐ Volume☐ Variety☐ Velocity☐ Variable 

26

27

28

29

30

31

32

33

34

35

<

31 of 40

>

Question # 31

UDF stands for:

 Revisit

Choose the best option

- ☐ Universal Defined Function
- ☐ Unique Defined Function
- ☐ Universal Disk Format
- ☐ Unique Definition of Function

Question # 32

 Revisit

Choose the best option

Can you provide multiple input paths to map reduce jobs?

- ☐ Yes developers can add any number of input paths
- ☐ Yes, but limit is currently 10 input paths
- ☐ No, hadoop works only on one input directory
- ☐ None of the above

Question # 26

 Revisit

The CREATE statement in Hive is related to:

Choose the best option

- ☐ DDL statements
- ☐ DML statements
- ☐ Session control statements
- ☐ Embedded SQL statements

Question # 22



Revisit

One of your HBase Region server (in a well configured in a sized HBase and Hadoop cluster) is reporting bad performance (slow response). What can be the possible reason?

Choose the best option

- ☐ small rows and column names
- ☐ uneven key space distribution
- ☐ small column family name
- ☐ None of the above



Question # 25

 Revisit

Choose the best option

Hadoop framework is written in:

- ☐ Java
- ☐ Python
- ☐ Scala
- ☐ C++

Question # 24

 Revisit

The parameter `fs.default.name` is set in _____ configuration file?

Choose the best option

- ☐ `hadoop-env.sh`
- ☐ `mapred-site.xml`
- ☐ `core-site.xml`
- ☐ `hdfs-site.xml`

Question # 18

 Revisit

Which of the following statements are correct?

Choose the best option

- ☐ Spark can run on the top of Hadoop
- ☐ Spark can process data stored in HDFS
- ☐ Spark can use Yarn as resource management layer
- ☐ All of the above

Question # 20

 Revisit

In YARN, a container can run:

Choose the best option

- ☐ any application or job developed to run on YARN
- ☐ only map and reduce tasks
- ☐ only java applications
- ☐ None of the above



6

7

8

9

10

11

12

13

14

15

<

14 of 40

>

All

13

27

Question # 14



What is default input format?

Choose the best option

- ☐ The default input format is xml. Developers can specify other input formats as appropriate if the xml is not correct input.
- ☐ There is no default input format. The input format should always be specified.
- ☐ The default input format is sequential input format. The data needs to be preprocessed before using the default input format.
- ☐ The default input format is Text input format with byte offset as key and entire line as value.

11

12

13

14

15

16

17

18

19

20

<

16 of 40

>

Question # 16

 Revisit

Which of the following are common feature of RDD and DataFrame?

Choose the best option

- ☐ immutability
- ☐ in-memory
- ☐ resilient
- ☐ All of the above

Question # 11

 Revisit

Which of the following deals with feature "structured but not relational"?

Choose the best option

- ☐ Structured data
- ☐ Unstructured Data
- ☐ Semi-structured data
- ☐ None of the above

Question # 10

 Revisit

Which of the following characteristic does not belong to big data?

Choose the best option

☐ Volume

☐ Variety

☐ Velocity

☐ Variable

6

7

8

9

10

11

12

13

14

15

<

13 of 40

>

All

Question # 13

 Revisit

Hive uses _____ to store metadata:

Choose the best option

- ☐ Derby database
- ☐ HiveQL
- ☐ NoSQL
- ☐ SQL

Question # 12

 Revisit

HBase is defined as _____

Choose the best option

- ☐ Row oriented
- ☐ Column oriented
- ☐ Tuple oriented
- ☐ None of the above



Question # 9

 Revisit

Choose the best option

Which of the following is true regarding apache airflow?

- ☐ open source
- ☐ workflow management platform
- ☐ data transformation pipeline ETL (Extract, Transform, Load)
- ☐ All of the above

Question # 1

 Revisit

Why are big data applications liable to latency?

Choose the best option:

- ☐ Big data cannot use in-memory computing.
- ☐ Big data applications are still in the early stages of development.
- ☐ The volume of big data is too large to be analyzed rapidly.
- ☐ Big data may reside in a different location from the application.